

Executive Summary

- **Taira sustained a 21.04% Moderate/High intent rate in Jul 22-28:** 235 qualified inquiries from 1,117 total. This was 0.87 percentage points below Jul 15-21, but 4.04 points above Jul 1-7.
- **Taira-origin inquiries converted better than the four selected Human CS in the current inquiry cohort:** 12 traceable bookings from 1,117 Taira inquiries (1.07%) versus 14 from 2,141 inquiries owned by Rem Reyes, Aira L. Garcia, Rolyn Ang, and Sarah (0.65%). This is a 1.64x current-week rate advantage, but the Jul 22-28 cohort is still maturing and Jul 28 is still partial.
- **Human CS booking conversion weakened in the latest cohort:** the combined booking rate moved from 1.29% in Jul 8-14 to 1.09% in Jul 15-21 and 0.65% in Jul 22-28. The latest period is less mature and Jul 28 is partial, but the decline is visible across all four agents and requires continued monitoring.
- **Sarah's weekly handoff service was mixed:** one-hour coverage improved from 73.68% to 79.04% and average business response time improved from 58.9 to 36.5 minutes, but the median slowed from 13.0 to 27.5 minutes and the 30-minute share fell from 62.41% to 51.53%. The setup improved parts of the distribution, not every SLA threshold.
- **Concurrent workload remains the strongest response-time constraint:** qualified-handoff median response previously rose from 11 minutes at 0-17 open cases to 134.5 minutes at 39+ open cases; one-hour coverage fell from 86.05% to 22.79%.
- **Guided selections reveal customer readiness:** Moderate/High intent increased from 12.93% with no guided selections to 25.38% with one and 61.54% with two. Contact capture remained near 5-6% at every depth and is the clearest funnel leak.

1. Chatbot Performance Snapshot

Bookings below are traceable to ManyChat inquiries and are attributed to the inquiry date, not the order creation date. There is no arbitrary conversion window; bookings observed through Jul 28 are included. The latest Jul 28 assignment snapshot was extracted at 22:38 Asia/Manila and remains partial.

Inquiry cohort	Inquiries	Moderate/High	Moderate rate	Bookings	Bookings / Moderate	Bookings / Inquiries
Jul 1-7	647	110	17.00%	9	8.18%	1.39%
Jul 8-14	934	183	19.59%	3	1.64%	0.32%
Jul 15-21	972	213	21.91%	13	6.10%	1.34%
Jul 22-28	1,117	235	21.04%	12	5.11%	1.07%

Jul 1-28	3,670	741	20.19%	37	4.99%	1.01%
-----------------	--------------	------------	---------------	-----------	--------------	--------------

Management readout

- Moderate/High intent quality remained above 20% in the latest two weeks.
- Weekly booking rates fluctuate because bookings often occur several days after the inquiry. Jul 22-28 should be revisited after additional maturation.
- The Jul 15-21 and Jul 22-28 booking cohorts are materially stronger than Jul 8-14, despite similar Moderate/High rates. This indicates that downstream selling activity and customer readiness, not intent classification alone, drive bookings.

2. Human CS Inquiry Conversion

Human CS is limited to **Rem Reyes, Aira L. Garcia, Rolyn Ang, and Sarah**. Each cell below shows **booking orders / inquiries (booking rate)**. Ownership is based on the original inquiry owner; booking creator is not used for conversion attribution.

Format and refresh: bookings/inquiries (conversion rate) . Jul 22-28 inquiry counts use the live original-assignment snapshot extracted Jul 28 at 22:38 Asia/Manila.

Agent	Jul 1-7	Jul 8-14	Jul 15-21	Jul 22-28	Jul 1-28
Aira L. Garcia	9/539 (1.67%)	8/647 (1.24%)	7/684 (1.02%)	5/666 (0.75%)	29/2,536 (1.14%)
Rem Reyes	6/540 (1.11%)	9/647 (1.39%)	7/678 (1.03%)	5/668 (0.75%)	27/2,533 (1.07%)
Rolyn Ang	3/325 (0.92%)	2/391 (0.51%)	7/463 (1.51%)	3/459 (0.65%)	15/1,638 (0.92%)
Sarah	5/541 (0.92%)	11/646 (1.70%)	3/376 (0.80%)	1/348 (0.29%)	20/1,911 (1.05%)
Four-CS team	23/1,945 (1.18%)	30/2,331 (1.29%)	24/2,201 (1.09%)	14/2,141 (0.65%)	91/8,618 (1.06%)

Management readout

- Aira led Jul 1-7; Sarah led Jul 8-14; Rolyn led Jul 15-21. No agent exceeded 0.75% in the still-maturing Jul 22-28 cohort.
- The latest cohort should be refreshed in subsequent reports because bookings created after Jul 28 will continue to accrue to these inquiry dates.

3. Sarah/Taira Qualified Handoff Performance

Response time is measured in **business minutes, 09:00-18:00 Asia/Manila**. For handoffs outside business hours, the clock begins at 09:00 the next business day rather than excluding the handoff.

Inquiry cohort	Qualified handoffs	Sarah responded	Response rate	Within 30 min	Within 1 hour	Median business response	Average business response	Customer replied after Sarah
Jul 15-21	133	130	97.74%	62.41%	73.68%	13.0 min	58.9 min	41.54%
Jul 22-28	229	224	97.82%	51.53%	79.04%	27.5 min	36.5 min	44.64%

Interpretation

- Response coverage remained essentially flat at about 98%.
- The latest week had fewer very long delays, reducing the average by 22.4 minutes and improving one-hour coverage by 5.36 points.
- The typical handoff was answered later: the median increased by 14.5 minutes and the 30-minute share fell by 10.88 points. The operational target should therefore be consistent early coverage, not only a lower average.
- In the clean Jul 25-26 full-Chatbot subset, the median improved from 28 to 23 business minutes versus Jul 20-23 and the 30-minute share increased from 50.65% to 57.89%, while one-hour coverage decreased from 80.52% to 78.07%. This supports better early prioritization for some leads, not a blanket SLA improvement.

Workload and commercial value

- With 0-17 open qualified cases, median response was 11 minutes and one-hour coverage was 86.05%. At 39+ open cases, median response increased to 134.5 minutes and one-hour coverage fell to 22.79%.
- Open queue and response delay had a Spearman correlation of 0.596 ($p < 0.001$). Queue size is the clearest operating control for response performance.
- In mature historical cohorts, Chatbot Moderate/High handoffs converted at 3.13% versus 0.84% for Sarah-direct inquiries. Median time-to-book was 18.0 hours for qualified handoffs versus 20.6 hours for Sarah-direct inquiries. This is evidence of lead-stage quality, not proof that the assignment setup alone caused the difference.

4. Guided Surfaces and Inquiry Progression

Cohorts cover **Jul 20-28**. A click identifies self-selected engagement and should not be interpreted as a randomized causal lift.

Surface	Shown	Clicked	Click rate	M/H: all shown	M/H: clicked	M/H: shown, no click	M/H: not shown
Promo carousel	103	67	65.05%	18.45%	20.90%	13.89%	18.40%
Price category	157	80	50.96%	30.77%	39.24%	22.08%	16.04%
Location buttons	269	103	38.29%	30.30%	60.19%	11.18%	13.38%
Product choices	224	19	8.48%	28.25%	47.37%	26.47%	14.64%

Behavioral signals

- **Location** is the strongest progression signal: 60.19% of clickers reached Moderate/High intent versus 11.18% of customers shown the buttons without clicking.
 - Location capture was 66.99% among clickers versus 10.56% among shown-no-click customers.
- Price-category clickers reached Moderate/High intent at 39.24%, 17.16 points above shown-no-click customers.
 - Choices skewed toward **Budget** (54.78%), followed by **Mid Range** (20.87%), **Premium** (18.26%), and **Economy** (6.09%).
- Product-choice engagement was low at 8.48%, but customers who clicked were more advanced in the journey.
 - Card-position clicks were concentrated on the first card (53.85%), followed by the second (26.92%).
- Promo action events were primarily Check Price (82.35%), then Promo Details (16.91%).
 - Promo brand actions were **Michelin** 42.86%, **Apollo** 36.61%, and **Yokohama** 20.54%.

Inquiry progression by guided selections

Moderate/High rates use conversations with sufficient analysis coverage; field rates use the inquiry cohort shown.

Guided selections	Inquiries	Moderate/High	Size captured	Brand captured	Location captured	Contact captured
-------------------	-----------	---------------	---------------	----------------	-------------------	------------------

0	696	12.93%	61.84%	40.03%	20.72%	5.92%
1	131	25.38%	92.31%	71.54%	33.85%	5.38%
2	52	61.54%	94.23%	80.77%	63.46%	5.77%
3+	7	57.14%	100.00%	85.71%	71.43%	0.00%

Actions

1. Make the next guided step explicit after every selection: size to price range, price range to product, and product to location or schedule.
2. Preserve previous selections and avoid asking customers to repeat captured information.
3. Move contact capture to a clear value exchange such as reserving stock, confirming branch availability, or sending the final quotation.
4. Keep the strongest default choices visible and instrument each impression, click, downstream field, handoff, and booking consistently.
5. Treat the 3+ selection cohort as directional only because it contains seven inquiries.

5. Conversation and Sales-Journey Insights

Human CS below includes only Rem Reyes, Aira L. Garcia, Rolyn Ang, and Sarah. The analysis covers Jul 22-28 conversations with semantic-analysis coverage.

Metric	Selected Human CS	Taira
Inquiries	2,141	1,117
Conversations analyzed	1,283 (59.93%)	1,010 (90.42%)
Semantic Moderate/High rate	15.12%	18.22%
Average exchanges	1.51	2.09
Reached pricing/quotation	25.33%	38.22%
Reached installation-partner stage	6.47%	9.21%
Reached scheduling stage	7.40%	10.20%
Captured contact information	5.53%	5.94%
Traceable inquiry-date booking rate	0.65%	1.07%

Interpretation

- Taira conversations had 38% more exchanges and were more likely to reach quotation, installation-partner, and scheduling stages. The current inquiry cohort also had a 1.64x booking rate. These are associated stage-depth signals; they do not establish that chatbot origin alone caused conversion.
- Both channels have nearly the same low contact-capture rate. The shared commercial opportunity is to convert completed product/location discovery into a low-friction contact or reservation step.
- Human CS analysis coverage is lower than Taira's. Coverage should be improved before using semantic rates as a complete productivity ranking.
- The most actionable near misses are conversations that capture size, brand, price preference, and location but stop before contact, stock confirmation, quotation, or scheduling.

6. GCP Cost Monitoring

Costs use the Cloud Billing export with the GCP UI's **Pacific charge-date basis**

(`America/Los_Angeles`) and net cost after credits. Gulong uses project `gulong-chatbot-459723` only.

Gulong

Service	Jul 1-7	Jul 8-14	Jul 15-21	Jul 22-28	Jul 1-28
Gemini API	\$91.25	\$73.98	\$74.93	\$84.72	\$324.88
BigQuery	\$48.10	\$101.41	\$24.87	\$37.64	\$212.02
Cloud Run	\$15.79	\$0.86	\$0.96	\$1.37	\$18.98
Compute Engine	\$20.41	\$18.42	\$16.95	\$13.01	\$68.79
Cloud SQL	\$3.05	\$3.16	\$3.53	\$3.20	\$12.94
Total	\$178.60	\$197.83	\$121.24	\$139.94	\$637.61

- Jul 22-28 cost increased 15.4% week over week, driven mainly by Gemini API and BigQuery, partly offset by lower Compute Engine cost.
- Straight-line Jul 1-28 projection for these five services is approximately **\$706** for July.

Mechanigo

Service	Jul 1-7	Jul 8-14	Jul 15-21	Jul 22-28	Jul 1-28
---------	---------	----------	-----------	-----------	----------

Gemini API	\$25.35	\$16.47	\$11.22	\$15.33	\$68.37
BigQuery	\$5.00	\$5.58	\$5.04	\$5.51	\$21.13
Cloud Run	\$1.11	\$0.73	\$0.08	\$0.06	\$1.98
Compute Engine	\$3.57	\$3.61	\$4.44	\$3.90	\$15.52
Cloud SQL	\$0.00	\$0.09	\$0.06	\$6.32	\$6.47
Total	\$35.03	\$26.48	\$20.84	\$31.12	\$113.47

- Jul 22-28 cost increased 49.3% week over week, mainly from Cloud SQL (+\$6.26) and Gemini API (+\$4.11).
- Straight-line July projection for these services is approximately **\$126**.

Carmax

Service	Jul 1-7	Jul 8-14	Jul 15-21	Jul 22-28	Jul 1-28
Gemini API	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00
BigQuery	\$1.49	\$1.79	\$1.79	\$1.58	\$6.65
Cloud Run	\$1.01	\$0.94	\$1.01	\$0.86	\$3.82
Compute Engine	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00
Cloud SQL	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00
Total	\$2.50	\$2.73	\$2.80	\$2.44	\$10.47

- Carmax remained stable and decreased 12.9% week over week.
- Straight-line July projection is approximately **\$12**.

7. Other Updates

RFID Inventory Warehouse

Updates

- Detection hit rate and scanning duration

Date Time	Actual Count	SKU	Scan Count	Scan time (s)	Hit %
7/22/2026 16:30:00	75	MAX POWER OIL FILTER MC-415	75	5	100.00%
7/23/2026 16:00:00	175	MAX POWER OIL FILTER MC-415	175	5.5	100.00%
7/25/2026 8:07	169	MAX POWER OIL FILTER MC-415	169	3.4	100.00%
7/27/2026 8:10	169	MAX POWER OIL FILTER MC-415	169	7.6	100.00%
7/28/2026 8:13	169	MAX POWER OIL FILTER MC-415	169	7.6	100.00%

Next steps

- Retest tagged inventory and document detection tolerance and hit rate.
- Compare internal printing against bulk supplier printing, including maintenance, ink, and printing-error costs.
- Evaluate the Android version for a custom bridge application that sends data to the internal system rather than directly to SIA.
- Pilot a phased rollout using oil filters, with a daily RFID count compared with SIA inventory.
- Consolidate supplier and operating documentation into an operations manual and use AI-assisted review to identify process optimizations.

▼ Reporting Definitions and Data Sources

- **Weekly periods:** Wednesday-Tuesday; Jul 1-7, Jul 8-14, Jul 15-21, and Jul 22-28.
- **Observation cutoff:** Jul 28, 2026; assignment snapshot extracted at 22:38 Asia/Manila. Jul 28 remains partial and recent inquiry cohorts are not fully mature.

- **Inquiry and Moderate/High metrics:**

`gulong_reporting.p_looker_agent_daily_conversion` and
`gulong_core.moderate_intent_sessions`.

- **Response timing:** `manychat_data.agent_assignment_events`,

`manychat_data.messages`, and `gulong_core.inquiry_assignments`; business hours
09:00-18:00 Asia/Manila.

- **Booking truth:** qualifying orders reconciled to `gulong_core.orders_booked`, with bounded
transcript-supported recovery for missing curated links.

- **Booking attribution:** original inquiry owner and inquiry date; no fixed conversion window. Order
creator is retained only as a closing-workload view.

- **Guided journey and conversation analysis:** ManyChat messages, guided-surface event logs,
inquiry evidence, and semantic conversation analysis tables.

- **Cost source:** Cloud Billing detailed export; Pacific charge date and net cost after credits.
